www.sciencemag.org/content/356/6334/183/suppl/DC1

Supplementary Materials for

Semantics derived automatically from language corpora contain human-like biases

Aylin Caliskan,* Joanna J. Bryson,* Arvind Narayanan*

*Corresponding author. Email: aylinc@princeton.edu (A.C.); jjb@alum.mit.edu (J.J.B.); arvindn@cs.princeton.edu (A.N.)

Published 14 April 2017, *Science* **356**, 183 (2017)
DOI: 10.1126/science.aal4230

This PDF file includes:

Materials and Methods

Cosine similarity. Given two vectors $\boldsymbol{x} = \langle x_1, x_2, \ldots, x_n \rangle$ and $\boldsymbol{y} = \langle y_1, y_2, \ldots, y_n \rangle$, their cosine similarity can be calculated as:

$$cos(\boldsymbol{x}, \boldsymbol{y}) = \frac{\Sigma_{i=1}^{n} x_i \cdot y_i}{\sqrt{\Sigma_{i=1}^{n} x_i^2}\ \sqrt{\Sigma_{i=1}^{n} y_i^2}}$$

In other words, it is the dot product of the vectors after they have been normalized to unit length.

Applying the Word Embedding Factual Association Test (WEFAT). Now we discuss in more detail how we apply WEFAT in two cases. The first is to test if occupation word vectors embed knowledge of the gender composition of the occupation in the real world. We use data released by the Bureau of Labor Statistics in which occupations are categorized hierarchically, and for each occupation the number of workers and percentage of women are given (some data is missing). The chief difficulty is that many occupation names are multi-word terms whereas the pre-trained word vectors that we use represent single words. Our strategy is to convert a multi-word term into a single word that represents a superset of the category (e.g., chemical engineer $\rightarrow$ engineer), and filter out occupations where this is not possible. The resulting words are listed in the following section.

Our second application of WEFAT is to test if androgynous names embed knowledge of how often the name is given to boys versus girls. We picked the most popular names in each 10% window of gender frequency based on 1990 U.S. Census data. Here again there is a difficulty: some names are also regular English words (e.g., *Will*). State-of-the-art word embeddings are not yet sophisticated enough to handle words with multiple senses or meanings; all usages are lumped into a single vector. To handle this, we algorithmically determine how "name-like" each vector is (by computing the distance of each vector to the centroid of all the name vectors), and eliminate the 20% of vectors that are least name-like.

Caveats about comparing WEAT to the IAT. In WEAT, much like the IAT, we do not compare two words. Many if not most words have multiple meanings, which makes pairwise measurements “noisy”. To control for this, we use small baskets of terms to represent a concept. In every case we use word baskets from previous psychological studies, typically from the same study we are replicating. We should note that distances / similarities of word embeddings lack any intuitive interpretation. But this poses no problem for us: our results and their import do not depend on attaching meaning to these distances.

While the IAT applies to individual human subjects, the embeddings of interest to us are derived from the *aggregate* writings of humans on the web. These corpora are generated in an uncontrolled fashion and are not representative of any one population. The IAT has been used to draw conclusions about populations by averaging individual results over samples. Our tests of word embeddings are loosely analogous to such population-level IATs.

Nevertheless, this difference precludes a direct numerical comparison between human biases measured by the IAT and biases in corpora measured by our methods. With word embeddings, there is no notion of test subjects. Roughly, it is as if we are able to measure the mean of the association strength over all the “subjects” who collectively created the corpus. But we have no way to observe variation between subjects or between trials. We do report p-values and effect sizes resulting from the use of multiple *words* in each category, but the meaning of these numbers is entirely different from those reported in IATs.

Text, figures, and legends

Replicating our results with other corpora and algorithms. We repeated all the *WEAT* and *WEFAT* analyses presented above using a different pre-trained embedding: word2vec on a Google News corpus (*3*). The embedding contains 3 million word vectors, and the corpus contains about 100 billion tokens, about an order of magnitude smaller than the Common Crawl corpus. Therefore the less common terms (especially names) in our lists occur infrequently in this corpus. This makes replication harder, as the co-occurrence statistics are "noisier". Yet in all WEATs except one, we observed statistically significant effects ($p < .05$) and large effect sizes. The lone exception is the pleasantness association of young vs. old people's names, a test which has a small number of target concepts and relatively low keyword frequencies. Table S1 summarizes the results.

Further, we found that the gender association strength of occupation words is highly correlated between the GloVe embedding and the word2vec embedding (Pearson $\rho = 0.88$; Spearman $\rho = 0.86$). In concurrent work, Bolukbasi et al. (*6*) compared the same two embeddings, using a different measure of the gender bias of occupation words, also finding a high correlation (Spearman $\rho = 0.81$).

Stereotypes reflected in statistical machine translation. One application where we can observe cultural stereotypes reflected is Statistical machine translation (SMT), a common natural language processing task. Translations to English from many gender-neutral languages such as Turkish lead to gender-stereotyped sentences. For example, Google Translate converts these Turkish sentences with gender-neutral pronouns: "O bir doktor. O bir hemşire." to these English sentences: "He is a doctor. She is a nurse." We see the same behavior for Finnish, Estonian, Hungarian, and Persian in place of Turkish. Similarly, translating the above two Turkish sentences into several of the most commonly spoken languages (Spanish, English, Portuguese, Russian, German, and French) results in gender-stereotyped pronouns in every case.

List of stimuli. Here we list the stimuli used in our WEAT and WEFAT tests. The WEAT tests are listed in the same order as Table 1.

WEAT 1: We use the flower and insect target words along with pleasant and unpleasant attributes found in (*5*).

- Flowers: aster, clover, hyacinth, marigold, poppy, azalea, crocus, iris, orchid, rose, bluebell, daffodil, lilac, pansy, tulip, buttercup, daisy, lily, peony, violet, carnation, gladiola, magnolia, petunia, zinnia.
- Insects: ant, caterpillar, flea, locust, spider, bedbug, centipede, fly, maggot, tarantula, bee, cockroach, gnat, mosquito, termite, beetle, cricket, hornet, moth, wasp, blackfly, dragonfly, horsefly, roach, weevil.
- Pleasant: caress, freedom, health, love, peace, cheer, friend, heaven, loyal, pleasure, diamond, gentle, honest, lucky, rainbow, diploma, gift, honor, miracle, sunrise, family, happy, laughter, paradise, vacation.
- Unpleasant: abuse, crash, filth, murder, sickness, accident, death, grief, poison, stink, assault, disaster, hatred, pollute, tragedy, divorce, jail, poverty, ugly, cancer, kill, rotten, vomit, agony, prison.

WEAT 2: We use the musical instruments and weapons target words along with pleasant and unpleasant attributes found in (*5*).

- Instruments: bagpipe, cello, guitar, lute, trombone, banjo, clarinet, harmonica, mandolin, trumpet, bassoon, drum, harp, oboe, tuba, bell, fiddle, harpsichord, piano, viola, bongo, flute, horn, saxophone, violin.
- Weapons: arrow, club, gun, missile, spear, axe, dagger, harpoon, pistol, sword, blade, dynamite, hatchet, rifle, tank, bomb, firearm, knife, shotgun, teargas, cannon, grenade, mace, slingshot, whip.

- Pleasant: As per previous experiment with insects and flowers.

- Unpleasant: As per previous experiment with insects and flowers.

WEAT 3: We use the European American and African American names along with pleasant and unpleasant attributes found in (*5*). Names that are marked with italics are excluded from our replication. In the case of African American names this was due to being to infrequent to occur in GloVe's Common Crawl corpus; in the case of European American names an equal number were deleted, chosen at random.

- European American names: Adam, *Chip*, Harry, Josh, Roger, Alan, Frank, *Ian*, Justin, Ryan, Andrew, *Fred*, Jack, Matthew, Stephen, Brad, Greg, *Jed*, Paul, *Todd*, *Brandon*, *Hank*, Jonathan, Peter, *Wilbur*, Amanda, Courtney, Heather, Melanie, *Sara*, *Amber*, *Crystal*, Katie, *Meredith*, *Shannon*, Betsy, *Donna*, Kristin, Nancy, Stephanie, *Bobbie-Sue*, Ellen, Lauren, *Peggy*, *Sue-Ellen*, Colleen, Emily, Megan, Rachel, *Wendy* (deleted names in italics).

- African American names: Alonzo, Jamel, *Lerone*, *Percell*, Theo, Alphonse, Jerome, Leroy, *Rasaan*, Torrance, Darnell, Lamar, Lionel, *Rashaun*, Tyree, Deion, Lamont, Malik, Terrence, Tyrone, *Everol*, Lavon, Marcellus, *Terryl*, Wardell, *Aiesha*, *Lashelle*, Nichelle, Shereen, *Temeka*, Ebony, Latisha, Shaniqua, *Tameisha*, *Teretha*, Jasmine, *Latonya*, *Shanise*, Tanisha, Tia, Lakisha, Latoya, *Sharise*, *Tashika*, Yolanda, *Lashandra*, Malika, *Shavonn*, *Tawanda*, Yvette (deleted names in italics).

- Pleasant: caress, freedom, health, love, peace, cheer, friend, heaven, loyal, pleasure, diamond, gentle, honest, lucky, rainbow, diploma, gift, honor, miracle, sunrise, family, happy, laughter, paradise, vacation.

- Unpleasant: abuse, crash, filth, murder, sickness, accident, death, grief, poison, stink, assault, disaster, hatred, pollute, tragedy, bomb, divorce, jail, poverty, ugly, cancer, evil, kill, rotten, vomit.

WEAT 4: We use the European American and African American names from (*7*), along with pleasant and unpleasant attributes found in (*5*).

- European American names: Brad, Brendan, Geoffrey, Greg, Brett, *Jay*, Matthew, Neil, Todd, Allison, Anne, Carrie, Emily, Jill, Laurie, *Kristen*, Meredith, Sarah (names in italics deleted in GloVe experiments).

- African American names: Darnell, Hakim, Jermaine, Kareem, Jamal, Leroy, Rasheed, *Tremayne*, Tyrone, Aisha, Ebony, Keisha, Kenya, *Latonya*, Lakisha, Latoya, Tamika, Tanisha (names in italics deleted in GloVe experiments).

- Pleasant: caress, freedom, health, love, peace, cheer, friend, heaven, loyal, pleasure, diamond, gentle, honest, lucky, rainbow, diploma, gift, honor, miracle, sunrise, family, happy, laughter, paradise, vacation.

- Unpleasant: abuse, crash, filth, murder, sickness, accident, death, grief, poison, stink, assault, disaster, hatred, pollute, tragedy, bomb, divorce, jail, poverty, ugly, cancer, evil, kill, rotten, vomit.

WEAT 5: We use the European American and African American names from (*7*), along with pleasant and unpleasant attributes found in (*9*).

- European American names: Brad, Brendan, Geoffrey, Greg, Brett, *Jay*, Matthew, Neil, Todd, Allison, Anne, Carrie, Emily, Jill, Laurie, *Kristen*, Meredith, Sarah (names in italics deleted in GloVe experiments).

- African American names: Darnell, Hakim, Jermaine, Kareem, Jamal, Leroy, Rasheed, *Tremayne*, Tyrone, Aisha, Ebony, Keisha, Kenya, *Latonya*, Lakisha, Latoya, Tamika, Tanisha (names in italics deleted in GloVe experiments).
- Pleasant: joy, love, peace, wonderful, pleasure, friend, laughter, happy.
- Unpleasant: agony, terrible, horrible, nasty, evil, war, awful, failure.

WEAT 6: We use the male and female names along with career and family attributes found in (*9*).

- Male names: John, Paul, Mike, Kevin, Steve, Greg, Jeff, Bill.
- Female names: Amy, Joan, Lisa, Sarah, Diana, Kate, Ann, Donna.
- Career: executive, management, professional, corporation, salary, office, business, career.
- Family: home, parents, children, family, cousins, marriage, wedding, relatives.

WEAT 7: We use the math and arts target words along with male and female attributes found in (*9*).

- Math: math, algebra, geometry, calculus, equations, computation, numbers, addition.
- Arts: poetry, art, dance, literature, novel, symphony, drama, sculpture.
- Male terms: male, man, boy, brother, he, him, his, son.
- Female terms: female, woman, girl, sister, she, her, hers, daughter.

WEAT 8: We use the science and arts target words along with male and female attributes found in (*10*).

- Science: science, technology, physics, chemistry, Einstein, NASA, experiment, astronomy.
- Arts: poetry, art, Shakespeare, dance, literature, novel, symphony, drama.
- Male terms: brother, father, uncle, grandfather, son, he, his, him.
- Female terms: sister, mother, aunt, grandmother, daughter, she, hers, her.

WEAT 9: We use the mental and physical disease target words along with uncontrollability and controllability attributes found in (*23*).

- Mental disease: sad, hopeless, gloomy, tearful, miserable, depressed.
- Physical disease: sick, illness, influenza, disease, virus, cancer.
- Temporary: impermanent, unstable, variable, fleeting, *short-term*, brief, occasional (word2vec experiments used short instead of short-term).
- Permanent: stable, always, constant, persistent, chronic, prolonged, forever.

WEAT 10: We use young and old people's names as target words along with pleasant and unpleasant attributes found in (*9*).

- Young people's names: Tiffany, Michelle, Cindy, Kristy, Brad, Eric, Joey, Billy.
- Old people's names: Ethel, Bernice, Gertrude, Agnes, Cecil, Wilbert, Mortimer, Edgar.
- Pleasant: joy, love, peace, wonderful, pleasure, friend, laughter, happy.
- Unpleasant: agony, terrible, horrible, nasty, evil, war, awful, failure.

WEFAT 1 (occupations): We use the gender stimuli found in (*9*) along with the occupation attributes we derived from Bureau of Labor Statistics.

- **Careers** : technician, accountant, supervisor, engineer, worker, educator, clerk, counselor, inspector, mechanic, manager, therapist, administrator, salesperson, receptionist, librarian, advisor, pharmacist, janitor, psychologist, physician, carpenter, nurse, investigator, bartender, specialist, electrician, officer, pathologist, teacher, lawyer, planner, practitioner, plumber, instructor, surgeon, veterinarian, paramedic, examiner, chemist, machinist, appraiser, nutritionist, architect, hairdresser, baker, programmer, paralegal, hygienist, scientist.
- **Female attributes**: female, woman, girl, sister, she, her, hers, daughter.
- **Male attributes**: male, man, boy, brother, he, him, his, son.

WEFAT 2 (androgynous names): We use the gender stimuli found in (*9*) along with the most popular androgynous names from 1990's public census data as targets.

- **Names** : Kelly, Tracy, Jamie, Jackie, Jesse, Courtney, Lynn, Taylor, Leslie, Shannon, Stacey, Jessie, Shawn, Stacy, Casey, Bobby, Terry, Lee, Ashley, Eddie, Chris, Jody, Pat, Carey, Willie, Morgan, Robbie, Joan, Alexis, Kris, Frankie, Bobbie, Dale, Robin, Billie, Adrian, Kim, Jaime, Jean, Francis, Marion, Dana, Rene, Johnnie, Jordan, Carmen, Ollie, Dominique, Jimmie, Shelby.
- **Female and Male attributes**: as per previous experiment on occupations.

Tables and legends

Target words	Attrib. words	Original Finding				Our Finding			
		Ref	**N**	**d**	**p**	$\mathbf{N_T}$	$\mathbf{N_A}$	**d**	**p**
Flowers vs insects	Pleasant vs unpleasant	(*5*)	32	1.35	10^{-8}	25×2	25×2	1.54	10^{-7}
Instruments vs weapons	Pleasant vs unpleasant	(*5*)	32	1.66	10^{-10}	25×2	25×2	1.63	10^{-8}
Eur.-American vs Afr.-American names	Pleasant vs unpleasant	(*5*)	26	1.17	10^{-5}	32×2	25×2	0.58	10^{-2}
Eur.-American vs Afr.-American names	Pleasant vs unpleasant	(*7*)	Not applicable			18×2	25×2	1.24	10^{-3}
Eur.-American vs Afr.-American names	Pleasant vs unpleasant from (*5*)	(*7*)	Not applicable			18×2	8 × 2	0.72	10^{-2}
Male vs female names	Career vs family	(*9*)	$39k$	0.72	10^{-2}	8 × 2	8 × 2	1.89	10^{-4}
Math vs arts	Male vs female terms	(*9*)	$28k$	0.82	$< 10^{-2}$	8 × 2	8 × 2	0.97	.027
Science vs arts	Male vs female terms	(*10*)	91	1.47	10^{-24}	8 × 2	8 × 2	1.24	10^{-2}
Mental vs physical disease	Temporary vs permanent	(*23*)	135	1.01	10^{-3}	6 × 2	7 × 2	1.30	.012
Young vs old people's names	Pleasant vs unpleasant	(*9*)	$43k$	1.42	$< 10^{-2}$	8 × 2	8 × 2	$-.08$	0.57

Table S1: Summary of Word Embedding Association Tests using word2vec embeddings trained on the Google News corpus. The rows and columns are as in Table 1. For certain tests, the number of WEAT target words here is different than in Table 1, because in each case, we delete words not found in the corresponding word embedding.

References

1. M. Stubbs, *Text and Corpus Analysis: Computer-Assisted Studies of Language and Culture* (Blackwell, Oxford, 1996).

2. J. A. Bullinaria, J. P. Levy, Extracting semantic representations from word co-occurrence statistics: A computational study. *Behav. Res. Methods* **39**, 510–526 (2007). doi:10.3758/BF03193020 Medline

3. T. Mikolov, J. Dean, Distributed representations of words and phrases and their compositionality. *Adv. Neural Inf. Process. Syst.* **2013**, 3111–3119 (2013).

4. C. M. Bishop, *Pattern Recognition and Machine Learning* (Springer, London, 2006).

5. A. G. Greenwald, D. E. McGhee, J. L. Schwartz, Measuring individual differences in implicit cognition: The implicit association test. *J. Pers. Soc. Psychol.* **74**, 1464–1480 (1998). doi:10.1037/0022-3514.74.6.1464 Medline

6. T. Bolukbasi, K.-W. Chang, J. Y. Zou, V. Saligrama, A. T. Kalai, Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. *Adv. Neural Inf. Process. Syst.* **2016**, 4349–4357 (2016).

7. M. Bertrand, S. Mullainathan, Are Emily and Greg more employable than Lakisha and Jamal? A field experiment on labor market discrimination. *Am. Econ. Rev.* **94**, 991–1013 (2004). doi:10.1257/0002828042002561

8. M. Bertrand, D. Chugh, S. Mullainathan, Implicit discrimination. *Am. Econ. Rev.* **95**, 94–98 (2005). doi:10.1257/000282805774670365

9. B. A. Nosek, M. Banaji, A. G. Greenwald, Harvesting implicit group attitudes and beliefs from a demonstration web site. *Group Dyn.* **6**, 101–115 (2002). doi:10.1037/1089-2699.6.1.101

10. B. A. Nosek, M. R. Banaji, A. G. Greenwald, Math = male, me = female, therefore math ≠ me. *J. Pers. Soc. Psychol.* **83**, 44–59 (2002). doi:10.1037/0022-3514.83.1.44 Medline

11. B. A. Nosek, F. L. Smyth, N. Sriram, N. M. Lindner, T. Devos, A. Ayala, Y. Bar-Anan, R. Bergh, H. Cai, K. Gonsalkorale, S. Kesebir, N. Maliszewski, F. Neto, E. Olli, J. Park, K. Schnabel, K. Shiomura, B. T. Tulbure, R. W. Wiers, M. Somogyi, N. Akrami, B. Ekehammar, M. Vianello, M. R. Banaji, A. G. Greenwald, National differences in gender-science stereotypes predict national sex differences in science and math achievement. *Proc. Natl. Acad. Sci. U.S.A.* **106**, 10593–10597 (2009). doi:10.1073/pnas.0809921106 Medline

12. P. D. Turney, P. Pantel, From frequency to meaning: Vector space models of semantics. *J. Artif. Intell. Res.* **37**, 141 (2010).

13. J. Pennington, R. Socher, C. D. Manning, GloVe: Global Vectors for Word Representation. *EMNLP* **14**, 1532–1543 (2014).

14. T. MacFarlane, Extracting semantics from the Enron corpus, University of Bath, Department of Computer Science Technical Report Series; CSBU-2013-08; http://opus.bath.ac.uk/37916/ (2013).

15. W. Lowe, S. McDonald, The direct route: Mediated priming in semantic space, *Proceedings of the Twenty-second Annual Conference of the Cognitive Science Society* (LEA, 2000), pp. 806–811.

16. M. Sahlgren, The distributional hypothesis. *Ital. J. Linguist.* **20**, 33 (2008).

17. G. Lupyan, The centrality of language in human cognition. *Lang. Learn.* **66**, 516–553 (2016). doi:10.1111/lang.12155

18. S. Barocas, A. D. Selbst, Big data's disparate impact. *Calif. Law Rev.* **104**, 2477899 (2014).

19. C. Dwork, M. Hardt, T. Pitassi, O. Reingold, R. Zemel, Fairness through awareness, *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference* (ACM, 2012), pp. 214–226.

20. M. Feldman, S. A. Friedler, J. Moeller, C. Scheidegger, S. Venkatasubramanian, Certifying and removing disparate impact, *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (ACM, 2015), pp. 259–268.

21. K. R. Thórisson, Integrated A.I. systems. *Minds Mach.* **17**, 11–25 (2007).

22. M. Hanheide, M. Göbelbecker, G. S. Horn, A. Pronobis, K. Sjöö, A. Aydemir, P. Jensfelt, C. Gretton, R. Dearden, M. Janicek, H. Zender, G.-J. Kruijff, N. Hawes, J. L. Wyatt, Robot task planning and explanation in open and uncertain worlds. *Artif. Intell.* **2015**, j.artint.2015.08.008 (2015). doi:10.1016/j.artint.2015.08.008

23. L. L. Monteith, J. W. Pettit, Implicit and explicit stigmatizing attitudes and stereotypes about depression. *J. Soc. Clin. Psychol.* **30**, 484–505 (2011). doi:10.1521/jscp.2011.30.5.484