

Language use shapes cultural norms: Large scale evidence from gender

Molly Lewis^{1,2} & Gary Lupyan¹

¹ University of Wisconsin-Madison

² University of Chicago

Author Note

Portions of this manuscript appeared in Lewis & Lupyan, 2018, Cog. Sci Proceedings.

Correspondence concerning this article should be addressed to Molly Lewis, . E-mail:
mollyllewis@gmail.com

Abstract

Keywords: cultural norms, implicit association task (IAT), gender

Word count: X

Language use shapes cultural norms: Large scale evidence from gender

Introduction

The language we use to communicate a message shapes how our listener interprets that message (Fausey & Boroditsky, 2010; Loftus & Palmer, 1974; Tversky & Kahneman, 1981). A listener, for example, is more likely to infer that a person is at fault if the event is described actively (e.g., “she ignited the napkin”), as opposed to passively (e.g., “the napkin ignited”). The formative power of language is perhaps most potent in shaping meanings that necessarily must be learned from others: cultural norms. In the present paper, we consider one type of cultural norm—gender—and examine the extent to which differences in language use may lead to cross-cultural differences in understandings of gender.

Gender provides a useful case study of the relationship between language and thought for several reasons. First, more abstract domains like gender may be more subject to the influence of language relative to more perceptually grounded domains like natural kinds (Boroditsky, 2001). Second, many languages encode the gender of speakers and addressees explicitly in their grammar. Third, a large body of evidence suggests that language plays a key role in transmitting social knowledge to children (e.g., Master, Markman, & Dweck, 2012). And, fourth, gender norms are highly variable across cultures and have clear and important social implications.

For our purposes, we define the hypothesis space of possible relationships between language and gender norms with two broad extremes: (1) language reflects a pre-existing gender bias in its speakers (*language-as-reflection hypothesis*); (2) language causally influences gender biases (*language-as-causal hypothesis*). We assume that the language-as-reflection hypothesis is true to some extent: some of the ways we talk about gender reflect our knowledge and biases acquired independently of language. For example, we may observe that most nurses are women, and therefore be more likely to use a female

pronoun to refer to a nurse of an unknown gender. Our goal here is to understand the extent to which language may also exert a causal influence on conceptualizations of gender.

In particular, we explore two possible mechanisms by which the way we speak may influence notions of gender¹. The first is through the overt grammatical marking of gender, particularly on nouns, which is obligatory in roughly one quarter of languages (e.g., in Spanish, “nina” (girl) and “enfermera” (nurse) both take the gender marker *-a* to indicate grammatical femininity; Corbett, 1991). Because grammatical gender has a natural link to the real world, speakers may assume that grammatical markers are meaningful even when applied to inanimate objects that do not have a biological sex. In addition, the mere presence of obligatory marking of grammatical gender may promote bias by making the dimension of gender more salient to speakers.

A second route by which language may shape gender norms is via word co-occurrences. Words that tend to occur in similar contexts in language may lead speakers to assume—either implicitly or explicitly—that they have similar meanings. For example, statistically, the word “nurse” occurs in many of the same contexts as the pronoun “her,” providing an implicit link between these two concepts that may lead to a bias to assume that nurses are female. This second route may be particularly influential because the bias is encoded in language in a way that is more implicit than grammatical markers of gender and thus more difficult to reject.

An existing body of experimental work points to a link between language and psychological gender bias in both adults (e.g., Phillips & Boroditsky, 2003) and children (e.g., Sera, Berge, & Castillo Pintado, 1994). For example, Phillips and Boroditsky (2003) asked Spanish-English and German-English adult bilinguals to make similarity judgements between pairs of pictures depicting an object with a natural gender (e.g., a bride) and one without (e.g., a toaster). They found that participants rated pairs as more similar when the pictures matched in grammatical gender in their native language. While these types of studies

¹These mechanisms are what Whorf (1945) refers to as phenotypes (overt) and cryptotypes (covert).

provide suggestive evidence for a causal link between language and psychological gender bias, they are limited by the fact that they typically only compare speakers of 2-3 different languages and measure bias in a way that is subject to demand characteristics.

In what follows, we ask whether the way gender is encoded linguistically across 31 different languages predicts cross-cultural variability in a particular manifestation of a gender bias—the bias to associate men with careers and women with family. We begin in Study 1 by describing cross-cultural variability in psychological gender bias using an implicit measure. In Study 1, we use semantic-embedding models to examine whether variability in lexical semantics predicts variability in psychological gender biases. In Study 2, we ask whether the presence of grammatical gender in a language is associated with greater implicit gender bias. Together, our data suggest that both language statistics and language structure likely play a causal role in shaping culturally-specific notions of gender.

Description of Cross-Cultural IAT data

To quantify cross-cultural gender bias, we used data from a large-scale administration of an Implicit Association Task (IAT; Greenwald, McGhee, & Schwartz, 1998) by Project Implicit (<https://implicit.harvard.edu/implicit/>; Nosek, Banaji, & Greenwald, 2002). The IAT measures the strength of respondents’ implicit associations between two pairs of concepts (e.g., male-career/female-family vs. male-family/female-career) accessed via words (e.g., “man,” “business”). The underlying assumption of the IAT is that words denoting more similar meanings should be easier to pair together compared to more dissimilar pairs.

Meanings are paired in the task by assigning them to the same response keys in a 2AFC categorization task. In the critical blocks of the task, meanings are assigned to keys in a way that is either bias-congruent (i.e. Key A = male/career; Key B = female/family) or bias-incongruent (i.e. Key A = male/family; Key B = female/career). Participants are then

presented with a word related to one of the four concepts and asked to classify it as quickly as possible. Slower reaction times in the bias-incongruent blocks relative to the bias-congruent blocks are interpreted as indicating an implicit association between the corresponding concepts (i.e. a bias to associate male with career and female with family).

In the present study, we analyzed a dataset of gender-career IAT scores collected by Project Implicit between 2005 and 2016. We restricted our sample based on participants' reaction times and error rates using the same criteria described in Nosek, Banaji, and Greenwald (2002, pg. 104). We only analyzed data for countries that had at least 400 participants. Our final sample included 769,353 participants from 40 countries, with a median of 1,302 participants per country. Note that although the respondents were from largely non-English speaking countries, the IAT was conducted in English. We do not have language background data from the participants, but we assume that most respondents from non-English speaking countries were native speakers of the dominant language of the country and L2 speakers of English.

Several measures have been used in the literature to quantify the strength of the bias from participants' responses on congruent and incongruent blocks on the IAT. Here, we used the most robust measure, D-score, which measures the difference between critical blocks for each participant while controlling for individual differences in response time (Greenwald, Nosek, & Banaji, 2003). In addition to the implicit measure, we also analyzed an explicit measure of gender bias. After completing the IAT, participants were asked, "How strongly do you associate the following with males and females?" for both the words "career" and "family." Participants indicated their response on a Likert scale ranging from *female* (1) to *male* (7). We calculated an explicit gender bias score for each participant as the Career response minus the Family response, such that greater values indicate a greater bias to associate males with career.

At the participant level, implicit bias scores were correlated with participant age such

Implicit psychological gender bias by country

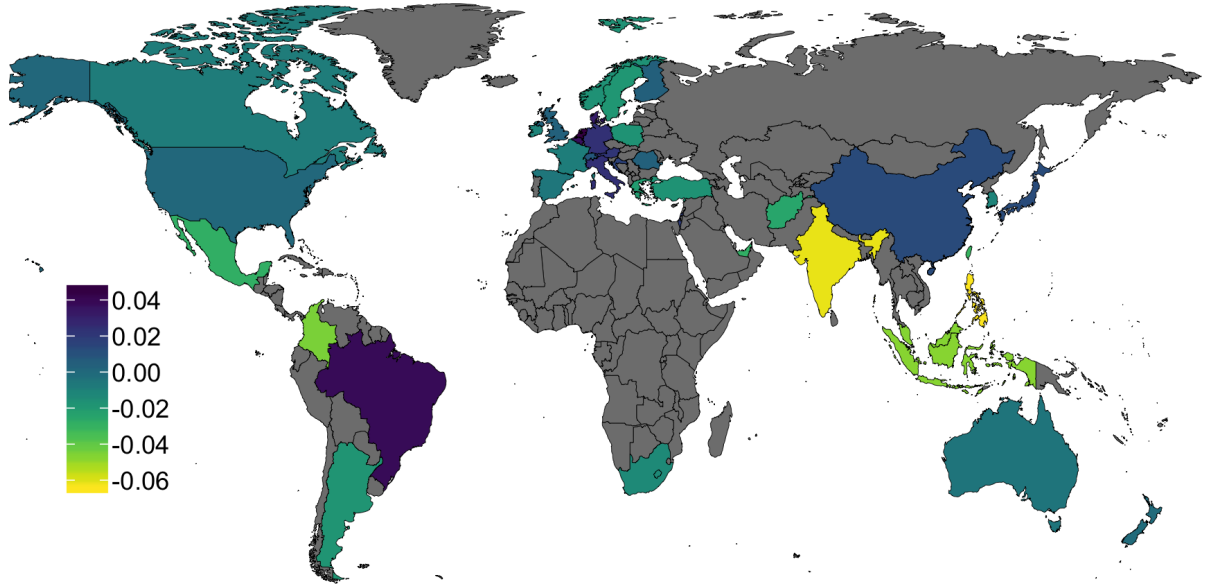


Figure 1. Residualized gender bias by country as measured by the IAT. Larger values indicate a larger bias to associate women with the concept of family and men with the concept of career. Countries in grey correspond to countries for which there was insufficient data to estimate the country-level gender bias.

that older participants tended to have a large gender bias compared to younger participants ($r = 0.06$, $p < .0001$). We also found that implicit bias scores varied as a function of participant sex with males ($M = 0.32$, $SD = 0.39$) showing a smaller implicit gender bias than females ($M = 0.42$, $SD = 0.36$; $t = 105.97$, $p < .0001$), a pattern consistent with previous findings (Nosek et al., 2002). Finally, implicit bias scores varied as a function of block order on the IAT task ($t = -114.59$, $p < .0001$). To account for these covariates, we calculated a residual implicit and residual explicit bias score for each participant, controlling for participant age, participant sex, and block order. We then averaged across participants to estimate the country-level gender bias. Figure 1 shows residualized implicit gender bias by country ($M = -0.01$; $SD = 0.03$). Implicit gender biases were moderately correlated with explicit gender biases at the level of participants ($r = 0.16$, $p < .0001$) but not countries ($r = 0.27$, $p = 0.09$).

We compared our residual country-level implicit and explicit gender biases to an objective measure of gender equality that is measured for each country by the United Nations Educational, Scientific and Cultural Organization (UNESCO): the percentage of women among science, technology, engineering, and mathematics (STEM) graduates in tertiary education (Miller, Eagly, & Linn, 2015; Stoet & Geary, 2018). We found that implicit gender bias was negatively correlated with percentage of women in STEM fields: Countries with a smaller gender bias tended to have more women in STEM fields ($r = -0.47$, $p = 0.01$). There was not a relationship between explicit gender bias and percentage of women in STEM fields ($r = 0.16$, $p = 0.43$). In addition, we found a strong correlation between the median age of a country, as measured by the CIA factbook (2017), and residual implicit bias: Countries with older populations tended to have larger gender biases ($r = 0.60$, $p < .0001$).

In sum, we replicate previously-reported patterns of gender bias in the gender-career IAT literature, with roughly comparable effect sizes (c.f. Nosek, et al., 2002). The weak correlation between explicit and implicit measures is consistent with claims that these two measures tap into different cognitive constructs (Forscher et al., 2016). In addition, consistent with previous research (Miller et al., 2015), we find that an objective measure of gender equality—female enrollment in STEM fields—is predictive of psychological gender bias.

Study 1: Gender bias and semantics

In Study 1, we ask whether participants’ implicit and explicit gender biases are correlated with the biases in the semantic structure of their native languages. For example, are the semantics of the words “woman” and “family” more similar in Hungarian than in English? Both the language-as-reflection and language-as-causal hypotheses predict a positive correlation between psychological and semantic gender biases.

As a model of word meanings, we use large-scale distributional semantics models

derived from auto-encoding neural networks trained on large corpora of text. The underlying assumption of these models is that the meaning of a word can be described by the words it tends to co-occur with—an approach known as distributional semantics (Firth, 1957). Under this approach, a word like “dog” is represented as more similar to “hound” than to “banana” because “dog” co-occurs with words more in common with “hound” than “banana.”

Recent developments in machine learning allow the idea of distributional semantics to be implemented in a way that takes into account many features of local language structure while remaining computationally tractable. The best known of these word embedding models is *word2vec* (Mikolov, Chen, Corrado, & Dean, 2013). The model takes as input a corpus of text and outputs a vector for each word corresponding to its semantics. From these vectors, we can derive a measure of the semantic similarity between two words by taking the distance between their vectors (e.g., cosine distance).

As it turns out, the biases previously reported using IAT tests can be predicted from distributional semantics models like *word2vec* using materials identical to those used in the IAT experiments. Caliskan, Bryson, and Narayanan (2017; henceforth *CBN*) measured the distance in vector space between the words presented to participants in the IAT task. CBN found that these distance measures were highly correlated with reaction times in the behavioral IAT task. For example, CBN find a bias to associate males with career and females with family in the career-gender IAT, suggesting that the biases measured by the IAT are also found in the lexical semantics of natural language.

CBN only measured semantic biases in English, however. In Study 1, we use the method described by CBN to measure gender bias in the range of first languages spoken by participants in the Project Implicit dataset by using models trained on those languages. In Study 1a, we first validate word embedding measures of gender bias by comparing them to explicit human judgements of gender bias. In Study 1b, we apply this method to models trained on text in other languages. We find that the implicit gender biases reported in Study

1 for individual countries are correlated with the biases found in the semantics of the natural language spoken by those participants.

Study 1a: Word embeddings as a measure of psychological gender bias

To validate word embeddings as a measure of psychological gender bias, we asked whether words that were closely associated with males in the word embedding models tended to be rated by human participants as being more male biased. We found human and word-embedding estimates of gender bias to be highly correlated.

Methods. We used an existing set of word norms in which participants were asked to rate “the gender associated with each word” on a Likert scale ranging from *very feminine* (1) to *very masculine* (7; Scott, Keitel, Becirspahic, Yao, & Sereno, 2018). We compared these gender norms to estimates of gender bias from a word embedding model pre-trained on the corpus of English Wikipedia using the fastText algorithm (a variant of word2vec; Bojanowski, Grave, Joulin, & Mikolov, 2016; Joulin, Grave, Bojanowski, & Mikolov, 2016). The model contains 2,519,370 words with each word represented by a 300 dimensional vector. To calculate a gender scores from the word embeddings, for each word we calculated the average cosine distance to a set of male words (“male”, “man”, “he”, “boy”, “his”, “him”, “son”, “brother”) and the average cosine similarity to a set of female words (“female”, “woman”, “she”, “girl”, “hers”, “her”, “daughter”, “sister”). A gender score for each word was then obtained by taking the difference of the similarity estimates (mean male similarity - mean female similarity), such that larger values indicated a stronger association with males. There were 4671 words in total that overlapped between the two data sources.

Results and Discussion. Estimates of gender bias from word embeddings ($M = 0$; $SD = 0.03$) and human judgements ($M = 4.10$; $SD = 0.92$) were highly correlated ($r = 0.59$; $p < .0001$; Fig. 2).




Figure 2. Word embedding estimates of gender bias as a function of human judgements of gender bias (Study 1a). Each point corresponds to a word. Larger numbers indicate stronger association with males. Blue line shows linear fit and the error band corresponds to a standard error (too small to be visible).

Study 1b: Cross-linguistic gender bias semantics

With our corpus validated, we next turn toward examining the relationship between psychological and linguistic gender biases. In Study 1b, we estimate the magnitude of the gender-career bias in each of the languages spoken in the countries described in the Project Implicit dataset and compare it with estimates of behavioral gender bias from the Project Implicit data set. If language plays a causal role in shaping psychological gender biases, we predicted that participants who speak a language with larger gender bias will tend to have a larger psychological gender bias.

Methods. For each country represented in our analysis of the Project Implicit, we identified the most frequently spoken language in each country using the CIA factbook (2017). This included a total of 27 unique languages. For a sample of 23 of these languages, we had native speakers translate a set of 8 female and 8 male target words into each language (identical to Study 1a). Native speakers also translated a set of 8 attributes words associated with the concept “career” (“career”, “executive”, “management”, “professional”, “corporation”, “salary”, “office”, “business”) and 8 attribute words associated with the concept “family” (“family”, “home”, “parents”, “children”, “cousins”, “marriage”, “wedding”, “relatives”). Our word lists are identical to the stimuli in the Project Implicit gender-career IAT behavioral task (Nosek et al., 2002) with one slight modification. In the behavioral task, proper names were used to cue the male and female categories (e.g. “John,” “Amy”). Because there are not direct translation equivalents of proper names, we instead used a set of generic gendered words which had been previously used for a different version of the gender IAT (e.g., “man,” “woman;” Nosek et al., 2002).

We used these translations to calculate a gender bias effect size from word embedding models trained on text in each language. Our effect size measure is a standardized difference score of the relative similarity of the target words to the target attributes (i.e. relative similarity of male to career vs. relative similarity of female to career). Our effect size measure is identical to the we used the the measure used in CBN with one exception (see SM for replication of CBN on our corpora). Namely, for languages with grammatically gendered attribute words (e.g., *ninas* for female children in Spanish), we calculated the relationship between target words and attribute words of the same gender (i.e. “man” to “*niños*” and “woman” to “*ninas*”). Our effect size measure is analogous to the to the behavioral effect size measure obtained from the IAT task in Project Implicit, where larger values indicate larger gender bias.

We calculated gender bias estimates based on models trained on two different corpora:

Wikipedia (Bojanowski et al., 2016) and Subtitles. These two models capture different types of XX. We then compared the effect size of the linguistic gender bias to the behavioral IAT gender bias estimated from Project Implicit, averaging across countries that speak whose participants speak the same language.




Figure 3. Residualized behavioral IAT gender bias as a function of semantic IAT gender bias by language (Study 1b). Semantic biases are estimated from models trained on each language using a subtitle corpus (left) and a sample of Wikipedia (right). Larger values indicate a larger bias to associate men with the concept of career and women with the concept of family. Error bands correspond to standard errors.

Results. At the level of languages, our estimate of gender bias in the semantics of each language was positively correlated the implicit gender bias of participants who spoke that language ($r = 0.51$; $p = 0.01$; Fig. 3). Importantly, this relationship held when controlling for median country age in a linear model ($\beta = 0.03$; $SE = 0.01$, $p = 0.03$). Semantic IAT gender bias was not correlated with explicit gender bias ($r = 0.14$; $p = 0.53$).

Table 1 shows the language-level correlations between all variables.

Table 1

Correlation (Pearson's r) for all measures in Study 1 at the level of languages. Asterisks indicate significance at the .05 level.

	Language IAT (Subtitles)	Language IAT (Wikipedia)	Residualized Explicit Bias	Residualized Behavioral IAT
Language IAT (Subtitles)				
Language IAT (Wikipedia)	.51*			
Residualized Explicit Bias	-.18	.14		
Residualized Behavioral IAT	.47	.51*	.10	
Perecent Women in Stem	-.43	-.29	.34	-.42

Discussion.

Study 2: Gender bias and grammar

The findings in Study 1 are consistent with both the language-as-causal and language-as-reflection hypotheses. In Study 2, we try to distinguish between the two hypotheses by asking whether there is a relationship between psychological gender bias and language along a linguistic dimension that is unlikely to be a subject of rapid change—namely, grammatical gender. While of course grammars do change over time, they are less malleable than the meanings of individual words, and thus less likely to be affected by psychological biases. We predict, therefore, that if language causally influences psychological gender biases, languages that encode gender grammatically will tend to have larger psychological gender biases.

Method. We coded each of the languages in our sample (Study 1) for grammatical gender. We used a coarse binary coding scheme, categorizing a language as encoding

grammatical gender if it made any gender distinction on noun classes (masculine, feminine, common or neuter), and as not encoding gender grammatically otherwise.

Results and Discussion.

General Discussion

References

- Bojanowski, P., Grave, E., Joulin, A., & Mikolov, T. (2016). Enriching word vectors with subword information.
- Boroditsky, L. (2001). Does language shape thought?: Mandarin and English speakers' conceptions of time. *Cognitive Psychology*, 43(1), 1–22.
- Caliskan, A., Bryson, J. J., & Narayanan, A. (2017). Semantics derived automatically from language corpora contain human-like biases. *Science*, 356(6334), 183–186.
- Central Intelligence Agency (CIA). (2017). The World Factbook. Retrieved from <https://www.cia.gov/library/publications/the-world-factbook/index.html>
- Corbett, G. G. (1991). *Gender*. Cambridge: Cambridge University Press.
- EF English Proficiency Index*. (2017). Retrieved from <https://www.ef.edu/epi/>
- Fausey, C. M., & Boroditsky, L. (2010). Subtle linguistic cues influence perceived blame and financial liability. *Psychonomic Bulletin & Review*, 17(5), 644–650.
- Firth, J. (1957). A synopsis of linguistic theory 1930-1955 in studies in linguistic analysis, philological society. Oxford.
- Forscher, P. S., Lai, C., Axt, J., Ebersole, C. R., Herman, M., Devine, P. G., & Nosek, B. A. (2016). A meta-analysis of change in implicit bias.
- Greenwald, A. G., McGhee, D. E., & Schwartz, J. L. (1998). Measuring individual differences in implicit cognition: The implicit association test. *Journal of Personality and Social Psychology*, 74(6), 1464.
- Greenwald, A. G., Nosek, B. A., & Banaji, M. R. (2003). Understanding and using the

- Implicit Association Test: An improved scoring algorithm. *Journal of Personality and Social Psychology*, 85(2), 197.
- Joulin, A., Grave, E., Bojanowski, P., & Mikolov, T. (2016). Bag of tricks for efficient text classification. *arXiv Preprint arXiv:1607.01759*.
- Loftus, E. F., & Palmer, J. C. (1974). Reconstruction of automobile destruction: An example of the interaction between language and memory. *Journal of Verbal Learning and Verbal Behavior*, 13(5), 585–589.
- Master, A., Markman, E., & Dweck, C. (2012). Thinking in categories or along a continuum: Consequences for children’s social judgments. *Child Development*, 83(4).
- Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Efficient estimation of word representations in vector space.
- Miller, D. I., Eagly, A. H., & Linn, M. C. (2015). Women’s representation in science predicts national gender-science stereotypes: Evidence from 66 nations. *Journal of Educational Psychology*, 107(3), 631.
- Nosek, B. A., Banaji, M. R., & Greenwald, A. G. (2002). Harvesting implicit group attitudes and beliefs from a demonstration web site. *Group Dynamics: Theory, Research, and Practice*, 6(1), 101.
- Phillips, W., & Boroditsky, L. (2003). Can quirks of grammar affect the way you think? Grammatical gender and object concepts. In *Proceedings of the 25th Annual Meeting of the Cognitive Science Society* (pp. 928–933).
- Scott, G. G., Keitel, A., Becirspahic, M., Yao, B., & Sereno, S. C. (2018). The glasgow norms: Ratings of 5,500 words on nine scales. *Behavior Research Methods*, 1–13.
- Sera, M. D., Berge, C. A., & Castillo Pintado, J. del. (1994). Grammatical and conceptual

- forces in the attribution of gender by English and Spanish speakers. *Cognitive Development*, 9(3), 261–292.
- Stoet, G., & Geary, D. C. (2018). The gender-equality paradox in science, technology, engineering, and mathematics education. *Psychological Science*, 29(4), 581–593.
- Tversky, A., & Kahneman, D. (1981). The framing of decisions and the psychology of choice. *Science*, 211(4481), 453–458.
- Whorf, B. (1945). Grammatical categories. *Language*, 1–11.
- Women, Peace and Security Index. (2017). Retrieved from <https://giwps.georgetown.edu/>